

Primer Aviso Cost Model

Cost Model

Build: Primer Aviso. Multi-agent first notice of loss for catastrophe surges.

This project has no recurring cost. There is no always-on resource, so it costs nothing when idle and bills only per case. The cost model below records each billable component, its pricing basis, and its estimated model calls per case.

Component | Pricing basis | Estimated cost per case

gpt-5-mini (Safety gate) | Per-call model usage | 0 to 1 call (if keyword check does not fire)

gpt-5-mini (Intake) | Per-call model usage | 3 to 5 calls (multi-turn conversation)

gpt-5-mini (Coverage) | Per-call model usage | 1 call (with one tool call to deterministic lookup)

gpt-5-mini (Evidence) | Per-call model usage | 1 call (multimodal, 1 to 2 images)

Fixed standing costs (vector store, reserved throughput, always-on endpoint) | None | \$0.00

Hosting | Not included | \$0.00 (orchestration runs locally)

Assumptions: single-user testing volume, gpt-5-mini deployed under Data Zone Standard EU in Sweden Central, and orchestration running locally.

What this means in practice: the total per case is roughly 6 to 8 gpt-5-mini calls. The only cost is per-call model usage. A hosted version would add a small, predictable compute cost (a container or function running intermittently, not continuously).

The cost lever

Within the per-call model usage, the prompt and response length controls the bill. Using short prompts and structured JSON output, rather than long-form text, keeps the output short and therefore keeps cost down.

Production cost shape

Not applicable to this build.