

Cost Model

Project 2: Startup Screening Workflow. Deterministic scoring on AWS Bedrock Flows.

Cindy Moncada | AI Solutions Architect

This project has no recurring cost. There is no always-on resource, so it costs nothing when idle and bills only per screening. The cost model below records each billable component, its pricing basis, and its estimated cost at single-user test volume, then identifies the design choices that control the bill.

Component	Pricing basis	Estimated cost at test volume
Claude Sonnet 4.5 (scoring)	\$3.00 / 1M input, \$15.00 / 1M output tokens	~\$0.015 per screening (~\$1.50 per 100)
Bedrock Guardrail (PII + content, in and out)	PII \$0.10, content \$0.15 per 1,000 text units	~\$0.001 per screening
Bedrock Flows orchestration	\$0.035 per 1,000 node transitions	Negligible at test volume
Storage / vector store	None required	\$0.00

Assumptions: single-user testing volume, a few thousand characters per submission, Claude Sonnet 4.5 on the Frankfurt regional endpoint, and on-demand pay-per-use pricing throughout. A guardrail text unit is up to 1,000 characters, and the guardrail evaluates both input and output, so each request is filtered twice.

The full ten-case evaluation harness cost roughly sixteen cents to run, and the screener costs nothing between runs because nothing stays provisioned. The model call is essentially the entire bill; the guardrail and the Flow orchestration are rounding errors at this scope. The 10 USD budget alert is watching a bill that, at test volume, will not come close to triggering it.

The cost lever

Within the model cost, output tokens cost five times more than input tokens on Sonnet 4.5, so the length of the assessment drives the bill more than the length of the startup description. The bounded, structured output format, a fixed score and one-line justification per dimension rather than open-ended prose, keeps the output short and therefore keeps cost down. This is a cost-conscious design decision, not only a formatting one.

One honest trade-off: choosing Frankfurt for EU residency carries roughly a 10% premium over the global endpoint for Sonnet 4.5 inference. That is a deliberate trade, residency over a marginal saving, and naming it shows the cost of the compliance posture was understood rather than stumbled into.

Production cost shape

This project has no fixed floor and scales linearly, roughly 16 USD per thousand screenings, with the bill falling to zero when idle. Project 1 was the opposite, dominated at low volume by a fixed, always-on search unit. The two builds have opposite cost shapes: one usage-metered with no floor, one floor-dominated until volume rises.